



InsureIntel AI

Agentic Insurance RAG & Regulatory Compliance Platform

MVP Business Documentation · v1.0 · 2025

CONFIDENTIAL — FOR BUSINESS AND INVESTOR USE ONLY

Prepared by: Abhishek and Gulshan | GenAI Engineering Lead

GitHub: github.com/alpha9934/insurance-ai-bot

Linkedin: <https://www.linkedin.com/in/abhishek-kumar-gupta-data-ai/> & <https://www.linkedin.com/in/gulshan-kumar-b07646355/>

1. Executive Summary

InsureIntel AI is a production-grade, enterprise Generative AI platform built to solve one of the most persistent operational challenges in the insurance industry: extracting accurate, actionable intelligence from massive, complex, and often encrypted policy documents — at speed, at scale, and without compliance risk.

Standard LLMs fail at this task. A 300-page insurance policy routinely exceeds 160,000 tokens, immediately colliding with API rate limits (30,000 TPM) and creating unacceptable hallucination risk on legally sensitive clauses. InsureIntel AI solves this through an Agentic RAG (Retrieval-Augmented Generation) architecture — acting as a precision-trained librarian that retrieves only the most relevant contractual context before querying the AI model.

Core Value Proposition: Answer any question about any insurance policy — accurately, instantly, and without exposing confidential documents to uncontrolled data pipelines.

FIELD	DETAIL
Product Name	InsureIntel AI
Version	MVP v1.0
Document Type	Business MVP Documentation
Classification	Confidential — Business & Investor Use
Primary Market	Enterprise Insurance (BFSI) — Underwriting, Claims, Compliance
Stage	Functional MVP — Production-Ready Architecture
License	Apache 2.0 (Open Source Core)
Repository	github.com/alpha9934/insurance-ai-bot

2. Problem Statement

The insurance industry processes millions of policy documents annually. Despite this volume, the intelligence locked inside these documents remains largely inaccessible to the people who need it most — adjusters, underwriters, compliance officers, and customer service teams.

2.1 The Three Critical Pain Points

Pain Point	Business Impact	Current Workaround
Document Overload	A single commercial policy can exceed 300 pages. Analysts spend 4–8 hours per policy review.	Manual review, keyword search — error-prone and slow.
Token Limit Collisions	LLMs cap at 30k TPM. Feeding a 160k-token document causes hard API failures (HTTP 429).	Document truncation — losing critical clauses.
Hallucination Risk	LLMs fabricate missing clause details. In insurance, a single wrong answer carries legal liability.	Human verification on every AI output — negates efficiency.
Encrypted Documents	Enterprise policies are AES-encrypted. Standard tools cannot even read them.	Manual decryption workflows, security gaps.
Regulatory Drift	State-level insurance laws change frequently. Policies referencing external regulations go stale.	Dedicated compliance teams — high headcount cost.

Market Signal: A 2024 McKinsey report estimates that AI-enabled claims processing and underwriting automation could unlock \$70–\$90 billion in annual productivity value globally across the insurance sector.

3. Solution Overview

InsureIntel AI addresses every identified pain point through a single, unified Agentic RAG pipeline. The platform is not a generic LLM wrapper — it is purpose-engineered for the structural and regulatory complexity of insurance documents.

3.1 How It Works — The Intelligence Pipeline

The system follows a five-stage processing loop:

- ▶ **Stage 1:** Document Ingestion & Decryption — AES-encrypted PDFs are natively decrypted at upload using pypdf + cryptography libraries. No manual pre-processing required.
- ▶ **Stage 2:** Semantic Chunking — The raw text is sliced into semantically coherent chunks (not arbitrary token splits). This preserves clause integrity and cross-reference relationships.
- ▶ **Stage 3:** Vector Embedding & Indexing — Each chunk is converted into a mathematical vector and stored in a local FAISS database. No cloud database subscription required for MVP.
- ▶ **Stage 4:** Context-Aware Retrieval — When a user submits a query, the engine performs a similarity search and retrieves the 5 most relevant clauses — and only those clauses.
- ▶ **Stage 5:** Agentic Answer Generation — GPT-4o receives the query + retrieved clauses. If the answer requires external regulatory data not in the document, the agent automatically calls the Tavily Search API to resolve the gap — without hallucinating.

3.2 Key Differentiators

Capability	InsureIntel AI	Standard LLM / RAG Tool
Encrypted PDF Support	Native AES decryption at ingestion	Not supported — fails at upload
Token Limit Handling	Chunked retrieval bypasses TPM limits entirely	Hard 429 errors on large docs
Hallucination Control	Retrieval-grounded + deterministic tool calling	Unconstrained generation
Regulatory Fact-Check	Autonomous Tavily web-search fallback agent	Static knowledge only
Data Privacy	In-memory processing, zero data egress	Data sent to cloud databases
Deployment	Fully containerized via Docker Compose	Complex infrastructure setup

4. Product Features — MVP Scope

4.1 Included in MVP v1.0

Feature	Description	Status
Encrypted PDF Upload	Drag-and-drop UI accepts AES-encrypted insurance PDFs. Decryption handled server-side.	✓ Live
Agentic Q&A Interface	Natural language chat interface. Users ask any policy question in plain English.	✓ Live
FAISS Local Vector DB	In-memory vector search. Lightning-fast, zero cost, zero egress risk.	✓ Live
Web Search Fallback	Tavily API integration — agent auto-searches when document context is insufficient.	✓ Live
GPT-4o Integration	LangChain orchestration with deterministic OpenAI function calling.	✓ Live
Containerized Deployment	Single docker compose up --build deploys entire stack.	✓ Live
Responsive UI	Next.js 15 + Tailwind CSS v4 frontend. Works on desktop and tablet.	✓ Live

4.2 Post-MVP Roadmap (v2.0+)

Feature	Business Value	Target Version
Persistent Vector Storage (Pinecone / pgvector)	Cross-policy querying — compare coverage across multiple contracts.	v2.0
Multi-Agent Orchestration (LangGraph)	Parallel agents for clause extraction + compliance cross-referencing.	v2.0
OCR / Vision Support (GPT-4o Vision)	Read scanned, handwritten, or image-heavy claim documents.	v2.1
Multi-Tenancy & Auth (Auth0 / Clerk)	Isolated workspaces per insurer / enterprise client.	v2.1
Audit Trail & Explainability	Every AI answer linked to its source clause. Regulator-ready output.	v2.2
IRDAI / State Compliance Module	Automated policy-vs-regulation gap analysis for Indian & US markets.	v3.0

5. Target Market & Use Cases

5.1 Primary Target Segments

SEGMENT	USE CASE
Insurance Carriers (Direct)	Internal tool for underwriting teams to query policy manuals and risk guidelines.
Third-Party Administrators (TPAs)	Claims adjusters need instant clause lookup without reading full contracts.
Insurance Brokers & Aggregators	Compare coverage across multiple carrier policies for a single client.
Legal & Compliance Teams	Scan new policy drafts against IRDAI, state, or NAIC mandates automatically.
InsurTech Platforms	Embed InsureIntel AI as the AI brain powering customer-facing policy chatbots.
Reinsurance Firms	Analyse treaty documents and exposure schedules at enterprise scale.

5.2 Illustrative User Stories

- ▶ **Adjuster:** A claims adjuster uploads a 240-page commercial property policy and asks: "Does Section 12 cover water damage from a burst pipe if the insured was not on premises?" — InsureIntel retrieves the exact sub-clause and delivers a legally-grounded answer in under 5 seconds.
- ▶ **Compliance:** A compliance officer asks: "Does this policy comply with Florida's new 2024 hurricane deductible disclosure requirements?" — The agent retrieves the relevant clause, then autonomously searches for the current Florida statute to verify alignment.
- ▶ **Underwriter:** An underwriter asks: "What is the aggregate liability limit for cyber incidents across all endorsements?" — The system scans all endorsement riders and synthesises the answer from multiple document sections.
- ▶ **Customer Service:** A policyholder support agent handles 3× more queries per hour because InsureIntel answers routine coverage questions instantly, escalating only edge cases to human specialists.

6. Technical Architecture Summary

This section provides a business-oriented summary of the technical design. Full technical documentation is available in the repository README.

6.1 Technology Stack

Layer	Technology	Business Rationale
Frontend	Next.js 15 (App Router) + Tailwind CSS v4	Modern, performant UI — fast to build, easy to white-label.
Backend API	FastAPI (Python) + Uvicorn	Asynchronous, high-throughput — handles concurrent document sessions.
AI Orchestration	LangChain + LangGraph	Industry standard for production agentic pipelines.
LLM	OpenAI GPT-4o	Best-in-class accuracy on legal/contractual language.
Vector Database	FAISS (local, in-memory)	Zero cost, zero egress for MVP. Swappable to Pinecone for v2.
Web Search Agent	Tavily AI Search API	Deterministic, LLM-optimised search — not a raw Google scrape.
Document Parsing	PyPDF + Cryptography	Handles AES-encrypted enterprise PDFs natively.
Containerisation	Docker + Docker Compose	One-command deployment — no environment config required.

6.2 Security & Data Handling

- ▶ All uploaded documents are processed entirely in-memory within the session. No policy content is written to disk or transmitted to external storage.
- ▶ Tavily API calls are governed by deterministic OpenAI function calling — the agent only calls external APIs when the document context is explicitly insufficient. Arbitrary data egress is architecturally prevented.
- ▶ API keys are injected at runtime via environment variables (.env). No credentials are hardcoded in the codebase.
- ▶ The Apache 2.0 licence allows enterprise clients to deploy the platform on-premises, ensuring documents never leave their own infrastructure.

7. Business Model & Monetisation

7.1 Proposed Revenue Streams

Model	Description	Target Customer
SaaS Subscription (Tier 1)	Hosted cloud platform. Per-seat pricing for teams of 5–50 analysts. ~\$200–500/seat/month.	Mid-size insurers, brokers, TPAs
Enterprise Licence (Tier 2)	On-premises deployment + source code licence. Annual contract. ~\$50k–200k/year.	Large carriers, reinsurers, GCCs
API Access (Tier 3)	Metered API for InsurTech platforms to embed policy Q&A into their own products.	InsurTech SaaS companies
Professional Services	Custom compliance module development, model fine-tuning on proprietary policy corpora.	Regulatory bodies, large carriers

7.2 Cost Structure (MVP)

- ▶ OpenAI API (GPT-4o): ~\$0.005–0.015 per query (retrieval-first architecture significantly reduces token consumption vs. naive LLM calls).
- ▶ Tavily API: Free tier sufficient for MVP; \$50/month for production volume.
- ▶ Infrastructure: ~\$50–200/month on a single cloud VM (AWS t3.medium or equivalent) for small-team deployments.
- ▶ FAISS (MVP): Zero cost. Pinecone migration (v2.0) adds ~\$70/month per namespace at scale.

Unit Economics: Estimated gross margin of 75–85% at SaaS scale, driven by low per-query AI costs and minimal infrastructure overhead for the retrieval-first architecture.

8. Competitive Landscape

Competitor	Strengths	InsureIntel AI Advantage
Generic RAG Platforms (e.g. LlamaIndex, Langchain Templates)	Flexible, open-source	Purpose-built for insurance: encrypted PDF support, Tavily compliance agent, BFSI use-case tuning
Shift Technology	Enterprise claims AI, strong brand	InsureIntel AI is open-core, deployable on-premises, no vendor lock-in
Hyperscaler AI (Azure OpenAI Service)	Scale and reliability	InsureIntel AI provides the insurance-specific orchestration layer that hyperscalers don't offer
Manual Review + Ctrl+F	No AI risk, zero cost	InsureIntel AI delivers 10–50x speed improvement with higher recall accuracy
ChatGPT / Copilot (consumer tools)	Familiar UX	No encrypted PDF support, no deterministic tool calling, no compliance audit trail

9. Risks & Mitigations

Risk	Likelihood / Impact	Mitigation
LLM Hallucination on Legal Clauses	Medium / High	Retrieval-grounded architecture limits generation to retrieved context. Audit trail in roadmap (v2.2).
OpenAI API Cost Escalation	Low / Medium	Retrieval-first design reduces tokens per query by ~80% vs naive prompting. Open-source LLM swap feasible.
Data Privacy / GDPR Compliance	Medium / High	In-memory processing, on-prem deployment option, no third-party data storage in MVP.
Regulatory Changes (IRDAI / NAIC)	Medium / Medium	Tavily web-search agent provides real-time regulatory context. Dedicated compliance module in v3.0 roadmap.
Competitor Replication	High / Medium	Insurance-domain tuning, encrypted PDF pipeline, and BFSI-specific agent tooling create meaningful moat.
Dependency on OpenAI	Medium / Medium	LangChain abstraction allows model swap to Anthropic Claude, Gemini, or open-source (Mistral, LLaMA).

10. MVP Success Metrics & KPIs

10.1 MVP Validation Metrics

Metric	Target (MVP)	Measurement Method
Query Accuracy (Retrieval)	>90% relevant clause retrieval on test policy set	Manual evaluation against 40-prompt benchmark suite
Average Response Time	<5 seconds per query (incl. vector search + GPT-4o)	Backend latency logging
429 Rate Limit Errors	0 errors on documents up to 200 pages	API error rate monitoring
Encrypted PDF Success Rate	100% decryption success on AES-256 test files	Upload pipeline testing
User Task Completion Rate	>85% of test users complete a query without assistance	Usability testing session

10.2 Commercial KPIs (Post-MVP)

- ▶ Time-to-first-customer: Target 1 signed pilot within 90 days of MVP demo.
- ▶ Analyst time saved per policy review: Target 60–70% reduction vs. manual baseline.
- ▶ Net Promoter Score (NPS): Target >40 among pilot users at 30 days.
- ▶ Monthly Recurring Revenue (MRR): Target ₹5L–10L (\$6k–12k) within 6 months of commercial launch.

11. Go-to-Market Strategy

Phase 1 — Validation (Months 1–3)

- ▶ Deploy InsureIntel AI as an open-source reference platform on GitHub. Build developer and InsurTech community visibility.
- ▶ Identify 2–3 pilot partners from BFSI/insurance contacts at Cognizant client network. Offer free pilot in exchange for structured feedback.
- ▶ Publish LinkedIn thought-leadership content demonstrating live policy Q&A demos. Target 5,000+ impressions per post.

Phase 2 — Pilot Commercialisation (Months 3–9)

- ▶ Convert pilot users to paid Tier 1 SaaS subscriptions. Price at ₹15k–25k/month for teams of 5–10 seats.
- ▶ Apply to InsurTech accelerators (e.g. NASSCOM DeepTech Club, Plug and Play InsurTech, Munich Re Ventures ecosystem).
- ▶ Develop 2–3 insurance-vertical case studies. Quantify analyst time saved, error rate reduction, compliance coverage improvement.

Phase 3 — Scale (Month 9+)

- ▶ Target GCC (Global Capability Centre) deals with India-based insurance back-office operations of global carriers.
- ▶ Launch enterprise on-premises offering for carriers with strict data sovereignty requirements.
- ▶ Explore strategic partnerships with LegalTech and RegTech platforms for embedded API distribution.

12. Appendix — Repository & Technical References

RESOURCE	LINK / DETAIL
GitHub Repository	github.com/alpha9934/insurance-ai-bot
Primary Language Split	TypeScript 49.5% · Python 37.4% · Dockerfile 5.8% · JS 4.4% · CSS 2.9%
Backend Entrypoint	backend/main.py — FastAPI + Uvicorn runtime
Frontend Entrypoint	frontend/src — Next.js 15 App Router + Tailwind CSS v4
AI Agent Module	backend/app/agents — LangChain OpenAI orchestration + Tavily tool binding
Container Config	docker-compose.yml — Two-service stack (backend :8000, frontend :3000)
Sample Policy	sample_policy.pdf — Test insurance document included in repository
Licence	Apache 2.0 — Commercial use permitted with attribution

— End of Document —